

Kubernetes hello-app

<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>

Deploying a containerized web application

This tutorial shows you how to package a web application in a Docker container image, and run that container image on a Google Kubernetes Engine cluster as a load-balanced set of replicas that can scale to the needs of your users.

Objectives

To package and deploy your application on GKE, you must:

1. Package your app into a Docker image
2. Run the container locally on your machine (optional)
3. Upload the image to a registry
4. Create a container cluster
5. Deploy your app to the cluster
6. Expose your app to the Internet
7. Scale up your deployment
8. Deploy a new version of your app

Build the container image

```
gcloud components install kubectl
```

```
gcloud config set project avian-compiler-238122
gcloud config set compute/zone us-central1-f
```

```
git clone https://github.com/GoogleCloudPlatform/kubernetes-engine-samples
cd kubernetes-engine-samples/hello-app
```

```
export PROJECT_ID=$(gcloud config get-value project -q)"
```

Upload the container image

```
docker build -t gcr.io/${PROJECT_ID}/hello-app:v1 .
gcloud auth configure-docker
docker push gcr.io/${PROJECT_ID}/hello-app:v1
```

The screenshot shows the Google Cloud Platform Container Registry interface. On the left, there's a sidebar with 'Container Registry' selected under 'Images'. The main area is titled 'My First Project' and lists a single repository entry: 'hello-app' with 'Hostname' set to 'gor.io'. The 'Visibility' is listed as 'Private'.

```
docker run --rm -p 8080:8080 gcr.io/${PROJECT_ID}/hello-app:v1
```

The terminal session shows the following steps:

- Building the Docker image: `docker build .` (output: `Successfully built 619b08ed77ee`)
- Tagging the image: `docker tag hello-app gcr.io/avian-compiler-238122/hello-app:v1`
- Pushing the image to GCR: `docker push gcr.io/PROJECT_ID/hello-app:v1`
- Running the Docker container locally: `docker run --rm -p 8080:8080 gcr.io/PROJECT_ID/hello-app:v1`
- Output: `2019/04/29 20:46:12 Server listening on port 8080`

Run your container locally (optional)

<https://8080-dot-7139517-dot-devshell.appspot.com/?authuser=0>



```
curl http://localhost:8080
```

Create a container cluster

```
gcloud container clusters create hello-cluster --num-nodes=2
```

```

creating cluster hello-cluster in us-central1-f... Cluster is being deployed...
Creating cluster hello-cluster in us-central1-f... Cluster is being health-checked (master is healthy)...done.
Created [https://container.googleapis.com/v1/projects/avian-compiler-238122/zones/us-central1-f/clusters/hello-cluster].
To inspect the contents of your cluster, go to: https://console.cloud.google.com/kubernetes/workload_gcloud/us-central1-f/hello-cluster?project=avian-compiler-238122
kubeconfig entry generated for hello-cluster.
NAME          LOCATION      MASTER VERSION   MASTER IP        MACHINE TYPE    NODE VERSION  NUM_NODES  STATUS
hello-cluster  us-central1-f  1.11.8-gke.6  104.154.178.255  n1-standard-1  1.11.8-gke.6  2          RUNNING
apigee@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$
apigee@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$

```

Kubernetes clusters

[CREATE CLUSTER](#) [DEPLOY](#) [REFRESH](#) [DELETE](#)

A Kubernetes cluster is a managed group of VM instances for running containerized applications. [Learn more](#)

<input type="checkbox"/> Name ^	Location	Cluster size	Total cores	Total memory	Notifications	Labels
<input checked="" type="checkbox"/> hello-cluster	us-central1-f	2	2 vCPUs	7.50 GB		Connect

Clusters

[EDIT](#) [DELETE](#) [DEPLOY](#) [CONNECT](#)

hello-cluster

- [Details](#)
- [Storage](#)
- [Nodes](#)

Cluster

Master version	1.11.8-gke.6	Upgrade available
Endpoint	104.154.178.255	Show credentials
Client certificate	Enabled	
Binary Authorization	Disabled	
Kubernetes alpha features	Disabled	
Total size	2	
Master zone	us-central1-f	
Node zones	us-central1-f	
Network	default	
Subnet	default	
VPC-native (alias IP)	Disabled	

The screenshot shows the 'Clusters' section of the Google Cloud Platform web interface. A cluster named 'hello-cluster' is selected. The 'Nodes' tab is active, displaying a table with two rows of node information. The columns are: Name, Status, CPU requested, CPU allocatable, Memory requested, Memory allocatable, Storage requested, and Storage allocatable. The first node is 'gke-hello-cluster-default-pool-bc51060a-49mx' and the second is 'gke-hello-cluster-default-pool-bc51060a-6rq7', both marked as 'Ready'.

gcloud compute instances list

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ gcloud compute instances list
NAME          ZONE      MACHINE_TYPE  PREEMPTIBLE INTERNAL_IP  EXTERNAL_IP
STATUS
gke-hello-cluster-default-pool-bc51060a-49mx  us-central1-f  n1-standard-1           10.128.0.17  35.238.173.95
RUNNING
gke-hello-cluster-default-pool-bc51060a-6rq7  us-central1-f  n1-standard-1           10.128.0.16  35.222.173.136
RUNNING
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122) $
```

gcloud container clusters get-credentials hello-cluster

★ Note: If you are using an existing Google Kubernetes Engine cluster or if you have created a cluster through Google Cloud Platform Console, you need to run the following command to retrieve cluster credentials and configure **kubectl** command-line tool with them:

```
gcloud container clusters get-credentials hello-cluster
```

If you have already created a cluster with the **gcloud container clusters create** command listed above, this step is not necessary.

Deploy your application

kubectl get pods

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pods
No resources found.
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122) $
```

```
kubectl run hello-web --image=gcr.io/${PROJECT_ID}/hello-app:v1 --port 8080
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
hello-web-6854fd845b-q4lps   1/1     Running   0          29s
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

Expose your application to the Internet

```
kubectl get service
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.19.240.1  <none>        443/TCP      10m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
kubectl expose deployment hello-web --type=LoadBalancer --port 80 --target-port 8080
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl expose deployment hello-web --type=LoadBalancer --port 80 --target-port 8080
service/hello-web exposed
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
hello-web  LoadBalancer  10.19.246.157  <pending>      80:31982/TCP  11s
kubernetes  ClusterIP  10.19.240.1  <none>        443/TCP      10m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
hello-web  LoadBalancer  10.19.246.157  <pending>      80:31982/TCP  22s
kubernetes  ClusterIP  10.19.240.1  <none>        443/TCP      11m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
hello-web  LoadBalancer  10.19.246.157  <pending>      80:31982/TCP  25s
kubernetes  ClusterIP  10.19.240.1  <none>        443/TCP      11m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
hello-web  LoadBalancer  10.19.246.157  <pending>      80:31982/TCP  34s
kubernetes  ClusterIP  10.19.240.1  <none>        443/TCP      11m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
hello-web  LoadBalancer  10.19.246.157  35.193.68.46  80:31982/TCP  56s
kubernetes  ClusterIP  10.19.240.1  <none>        443/TCP      11m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

<http://35.193.68.46/>

```
← → C ⓘ Not Secure | 35.193.68.46
```

```
Hello, norridge v1!
Version: 1.0.0
Hostname: hello-web-6854fd845b-q4lps
```

Scale up your application

```
kubectl scale deployment hello-web --replicas=3
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl scale deployment hello-web --replicas=3
deployment.extensions/hello-web scaled
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
kubectl get deployment hello-web
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get deployment hello-web
NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hello-web   3         3         3           3           7m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
kubectl get pods
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
hello-web-6854fd845b-q4lps   1/1     Running   0          8m
hello-web-6854fd845b-rcqs8   1/1     Running   0          1m
hello-web-6854fd845b-xd779   1/1     Running   0          1m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
hello-web-6854fd845b-q4lps   1/1     Running   0          8m
hello-web-6854fd845b-rcqs8   1/1     Running   0          1m
hello-web-6854fd845b-xd779   1/1     Running   0          1m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, norridge v1!
Version: 1.0.0
Hostname: hello-web-6854fd845b-q4lps
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, norridge v1!
Version: 1.0.0
Hostname: hello-web-6854fd845b-rcqs8
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, norridge v1!
Version: 1.0.0
Hostname: hello-web-6854fd845b-xd779
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, norridge v1!
Version: 1.0.0
Hostname: hello-web-6854fd845b-q4lps
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$
```

Deploy a new version of your app

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ vi main.go
```

```
// [START all]
package main

import (
    "fmt"
    "log"
    "net/http"
    "os"
)

func main() {
    // use PORT environment variable, or default to 8080
    port := "8080"
    if fromEnv := os.Getenv("PORT"); fromEnv != "" {
        port = fromEnv
    }

    // register hello function to handle all requests
    server := http.NewServeMux()
    server.HandleFunc("/", hello)

    // start the web server on port and accept requests
    log.Printf("Server listening on port %s", port)
    err := http.ListenAndServe(": "+port, server)
    log.Fatal(err)
}

// hello responds to the request with a plain-text "Hello, world" message.
func hello(w http.ResponseWriter, r *http.Request) {
    log.Printf("Serving request: %s", r.URL.Path)
    host, _ := os.Hostname()
    fmt.Fprintf(w, "Hello, hoffman v2!\n")
    fmt.Fprintf(w, "Version: 1.0.0\n")
    fmt.Fprintf(w, "Hostname: %s\n", host)
}
// [END all]
"main.go" 52L, 1435C
```

```
docker build -t gcr.io/${PROJECT_ID}/hello-app:v2 .
```

```
docker push gcr.io/${PROJECT_ID}/hello-app:v2
```

```

apiigefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ docker build -t gcr.io/${PROJECT_ID}/hello-app:v2 .
Sending build context to Docker daemon 9.728kB
Step 1/7 : FROM golang:1.8-alpine
--> 4cb86d3661bf
Step 2/7 : ADD ./go/src/hello-app
--> 68aa0a3ccca7
Step 3/7 : RUN go install hello-app
--> Running in 62bb1ecf34b5
Removing intermediate container 62bb1ecf34b5
--> 0dbc7b5e6637
Step 4/7 : FROM alpine:latest
--> cdf98d1859c1
Step 5/7 : COPY --from=0 /go/bin/hello-app .
--> 8e1b4ela1244
Step 6/7 : ENV PORT 8080
--> Running in de7f80e5ae56
Removing intermediate container de7f80e5ae56
--> 25e92a982982
Step 7/7 : CMD ["./hello-app"]
--> Running in 7f2d3bf2d1c4
Removing intermediate container 7f2d3bf2d1c4
--> d6d551b4cfcc3
Successfully built d6d551b4cfcc3
Successfully tagged gcr.io/avian-compiler-238122/hello-app:v2
apiigefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$
apiigefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ docker push gcr.io/${PROJECT_ID}/hello-app:v2
The push refers to repository [gcr.io/avian-compiler-238122/hello-app]
d7ce4055324: Pushed
a464c54f93a9: Layer already exists
v2: digest: sha256:f1ec51a539e2cb814deda5b6ea5fc115ad9bb71d3a2261a8ae663fe046e68a2a size: 739

```

kubectl get deployment hello-web
 kubectl get pods

```

apiigefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl set image deployment/hello-web hello-web=gcr.io/${PROJECT_ID}/hello-app:v2
deployment.extensions/hello-web image updated
apiigefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get deployment hello-web
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hello-web     3         3         3           3           19m
apiigefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
hello-web-56c4d8f675-jwhlm  1/1    Running   0          24s
hello-web-56c4d8f675-wqr27  1/1    Running   0          29s
hello-web-56c4d8f675-xhbkz  1/1    Running   0          26s
apiigefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$

```

<http://35.193.68.46/>

← → ⌂ ⓘ Not Secure | 35.193.68.46

Hello, hoffman v2!
 Version: 1.0.0
 Hostname: hello-web-56c4d8f675-jwhlm

```

apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
hello-web-56c4d8f675-jwhlm  1/1     Running   0          2m
hello-web-56c4d8f675-wqr27  1/1     Running   0          2m
hello-web-56c4d8f675-xhbkz  1/1     Running   0          2m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, hoffman v2!
Version: 1.0.0
Hostname: hello-web-56c4d8f675-wqr27
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, hoffman v2!
Version: 1.0.0
Hostname: hello-web-56c4d8f675-xhbkz
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, hoffman v2!
Version: 1.0.0
Hostname: hello-web-56c4d8f675-jwhlm
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, hoffman v2!
Version: 1.0.0
Hostname: hello-web-56c4d8f675-jwhlm
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ curl http://35.193.68.46/
Hello, hoffman v2!
Version: 1.0.0
Hostname: hello-web-56c4d8f675-wqr27
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$

```

Setting up HTTP Load Balancing with Ingress (This works as Load Balancer takes 10-15 minutes to propagate)

<https://cloud.google.com/kubernetes-engine/docs/tutorials/http-balancer>

kubectl get pod

```

apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pod
NAME           READY   STATUS    RESTARTS   AGE
hello-web-56c4d8f675-jwhlm  1/1     Running   0          11m
hello-web-56c4d8f675-wqr27  1/1     Running   0          11m
hello-web-56c4d8f675-xhbkz  1/1     Running   0          11m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$

```

Deploy a web application

```

kubectl run web1 --image=gcr.io/${PROJECT_ID}/hello-app:v1 --port 8080
kubectl run web2 --image=gcr.io/${PROJECT_ID}/hello-app:v2 --port 8080

```

kubectl get pod

```

apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get pod
NAME           READY   STATUS    RESTARTS   AGE
hello-web-56c4d8f675-jwhlm  1/1     Running   0          12m
hello-web-56c4d8f675-wqr27  1/1     Running   0          12m
hello-web-56c4d8f675-xhbkz  1/1     Running   0          12m
web1-5ffc78664b-hkphs      1/1     Running   0          18s
web2-976d7685c-jscst       1/1     Running   0          5s
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$

```

Expose your Deployment as a Service internally

```
kubectl expose deployment web1 --target-port=8080 --type=NodePort
kubectl expose deployment web2 —target-port=8080 --type=NodePort
```

```
kubectl get service web1
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service web
No resources found.
Error from server (NotFound): services "web" not found
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service web1
No resources found.
Error from server (NotFound): services "web1" not found
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl expose deployment web1 --target-port=8080 --type=NodePort
service/web1 exposed
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl expose deployment web2 -target-port=8080 --type=NodePort
service/web2 exposed
Error from server (NotFound): deployments.extensions "-target-port=8080" not found
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service web
No resources found.
Error from server (NotFound): services "web" not found
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service web1
NAME      TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)      AGE
web1     NodePort    10.19.246.86 <none>        8080:32332/TCP   1m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
kubectl get service web1
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service web1
NAME      TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)      AGE
web1     NodePort    10.19.246.86 <none>        8080:32332/TCP   2m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
vi basic-ingress.yaml
```

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: basic-ingress
spec:
  backend:
    serviceName: web1
    servicePort: 8080
~
```

```
kubectl apply -f basic-ingress.yaml
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl apply -f basic-ingress.yaml
ingress.extensions/basic-ingress created
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
kubectl get ingress basic-ingress
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get ingress basic-ingress
NAME      HOSTS      ADDRESS      PORTS      AGE
basic-ingress *          80          34s
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get ingress basic-ingress
NAME      HOSTS      ADDRESS      PORTS      AGE
basic-ingress *          34.95.110.35  80          1m
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

Setting up HTTP Load Balancing with Ingress - Part 2

Makes use of yaml

gcloud container clusters create loadbalancedcluster

<https://cloud.google.com/kubernetes-engine/docs/tutorials/http-balancer>

```
kubeconfig entry generated for loadbalancedcluster.
NAME      LOCATION      MASTER VERSION      MASTER_IP      MACHINE_TYPE      NODE_VERSION      NUM_NODES      STATUS
loadbalancedcluster us-central1-f 1.11.8-gke.6  104.154.73.174 n1-standard-1 1.11.8-gke.6 3      RUNNING
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl run web --image=gcr.io/google-samples/hello-app:1.0 --port=8080
deployment.apps/web created
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl expose deployment web --target-port=8080 --type=NodePort
service/web exposed
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get service web
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
web      NodePort      10.43.250.32 <none>      8080:31650/TCP      10s
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ vi basic-ingress.yaml
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl apply -f basic-ingress.yaml
ingress.extensions/basic-ingress created
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get ingress basic-ingress
NAME      HOSTS      ADDRESS      PORTS      AGE
basic-ingress *          80          9s
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

```
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get ingress basic-ingress
NAME      HOSTS      ADDRESS      PORTS      AGE
basic-ingress *          34.95.110.35  80          46s
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ █
```

★ **Note:** It may take a few minutes for GKE to allocate an external IP address and set up forwarding rules until the load balancer is ready to serve your application. In the meanwhile, you may get errors such as HTTP 404 or HTTP 500 until the load balancer configuration is propagated across the globe.

Google Cloud Platform My First Project

Network services Load balancing + CREATE LOAD BALANCER ⌂ REFRESH ⌁ DELETE

Load balancers Backends Frontends

Filter by name or protocol

Name	Protocol	Backends
k8s-um-default-basic-ingress-5c7997fc0ce705e2	HTTP	1 backend service (1 instance group, 0 network endpoint groups)
a0a3b1ded6ac211e988d942010a80016	TCP	1 target pool (2 instances)

To edit load balancing resources like forwarding rules and target proxies, go to the [advanced menu](#).

Not Secure | 34.95.110.35

Error: Server Error

The server encountered a temporary error and could not complete your request.

Please try again in 30 seconds.

Not Secure | 35.193.68.46

Hello, hoffman v2!
Version: 1.0.0
Hostname: hello-web-56c4d8f675-wqr27

[Load balancer details](#) [EDIT](#) [DELETE](#)

a0a3b1ded6ac211e988d942010a80016

Frontend

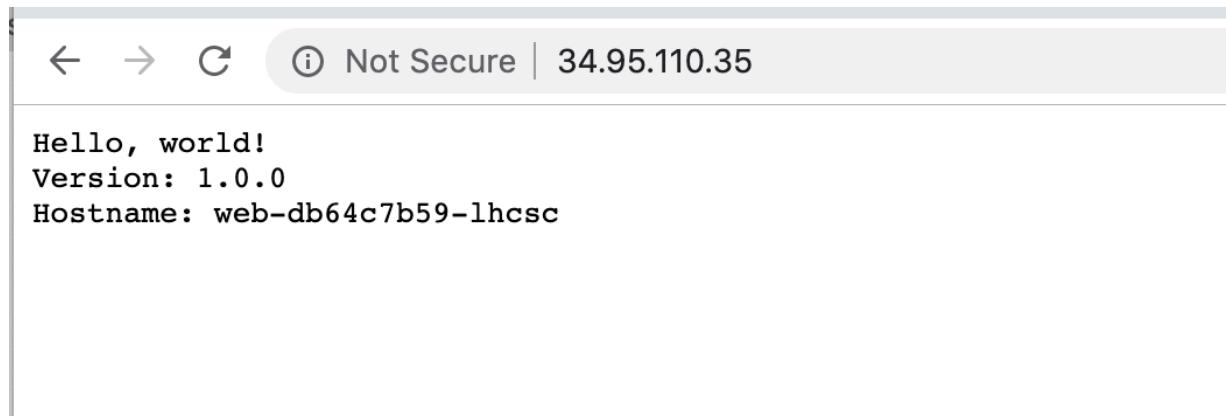
Protocol	IP:Port	Network Tier
TCP	35.193.68.46:80	Premium

Backend

Name: a0a3b1ded6ac211e988d942010a80016 Region: us-central1 Session affinity: None Health check: k8s-34a6f19edd6794c8-node

Instances	IP
gke-hello-cluster-default-pool-bc51060a-49mx	35.193.68.46
gke-hello-cluster-default-pool-bc51060a-6rq7	35.193.68.46

Ingress <http://34.95.110.35/>



kubectl run web2 --image=gcr.io/google-samples/hello-app:2.0 --port=8080

```

apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: fanout-ingress
spec:
  rules:
  - http:
    paths:
    - path: /*
      backend:
        serviceName: web
        servicePort: 8080
    - path: /v2/*
      backend:
        serviceName: web2
        servicePort: 8080

```

~

```

apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl run web2 --image=gcr.io/google-samples/hello-app:2.0 --port=8080
deployment.apps/web2 created
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl expose deployment web2 --target-port=8080 --type=NodePort
service/web2 exposed
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ vi fanout-ingress.yaml
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ vi fanout-ingress.yaml
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl create -f fanout-ingress.yaml
ingress.extensions/fanout-ingress created
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$

```

kubectl get ingress fanout-ingress

```

apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$ kubectl get ingress fanout-ingress
NAME          HOSTS   ADDRESS      PORTS   AGE
fanout-ingress *        34.95.99.222  80      47s
apigeefellow@cloudshell:~/kubernetes-engine-samples/hello-app (avian-compiler-238122)$

```

The screenshot shows the Google Cloud Platform Load Balancing interface. At the top, there are buttons for 'CREATE LOAD BALANCER', 'REFRESH', and 'DELETE'. Below this, a navigation bar has 'Load balancers' selected, with 'Backends' and 'Frontends' as other options. A search bar at the top right contains the placeholder 'Filter by name or protocol'. The main table lists three load balancers:

Name	Protocol	Backends
k8s-um-default-basic-ingress-5c7997fc0ce705e2	HTTP	1 backend service (1 instance group, 0 network endpoint groups)
k8s-um-default-fanout-ingress-5c7997fc0ce705e2	HTTP	3 backend services (1 instance group, 0 network endpoint groups)
a0a3b1ded6ac211e988d942010a80016	TCP	1 target pool (2 instances)

At the bottom, a note says: 'To edit load balancing resources like forwarding rules and target proxies, go to the advanced menu.'

k8s-um-default-fanout-ingress--5c7997fc0ce705e2

Details Monitoring Caching

Frontend

Protocol ^	IP:Port	Network Tier ?
HTTP	34.95.99.222:80	Premium

Host and path rules

Hosts ^	Paths	Backend
All unmatched (default)	All unmatched (default)	k8s-be-31148--5c7997fc0ce705e2
*	/*	k8s-be-31148--5c7997fc0ce705e2
*	/*	k8s-be-31650--5c7997fc0ce705e2
*	/v2/*	k8s-be-31345--5c7997fc0ce705e2

Backend

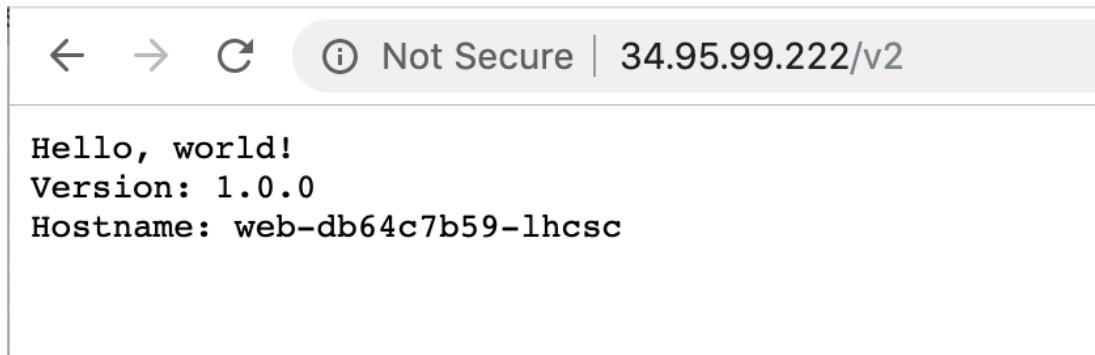
Backend services

1. k8s-be-31148--5c7997fc0ce705e2

Endpoint protocol: HTTP Named port: port31148 Timeout: 30 seconds Cloud CDN: disabled Health check: k8s-be-31148--5c7997fc0ce705e2

V2

<http://34.95.99.222/v2>



<http://34.95.99.222/>

← → C

i Not Secure | 34.95.99.222

Hello, world!

Version: 1.0.0

Hostname: web-db64c7b59-lhcsc